6. CPU Scheduling

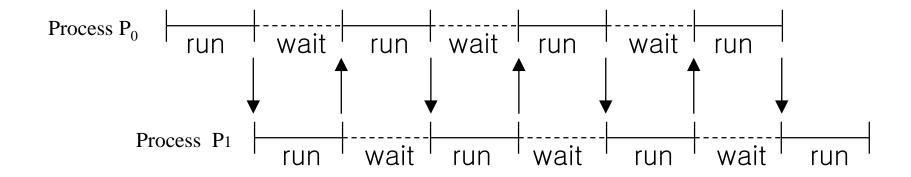
ECE30021/ITP30002 Operating Systems

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

Basic Concepts

- Motivation: maximum CPU utilization obtained with multiprogramming and multitasking
 - Resources (including CPU) are shared among processes



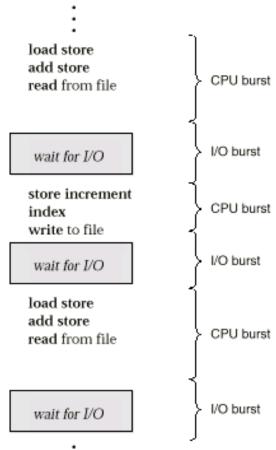
Multiprogramming

CPU-I/O Burst Cycle

Process execution consists of cycle of CPU execution and I/O wait

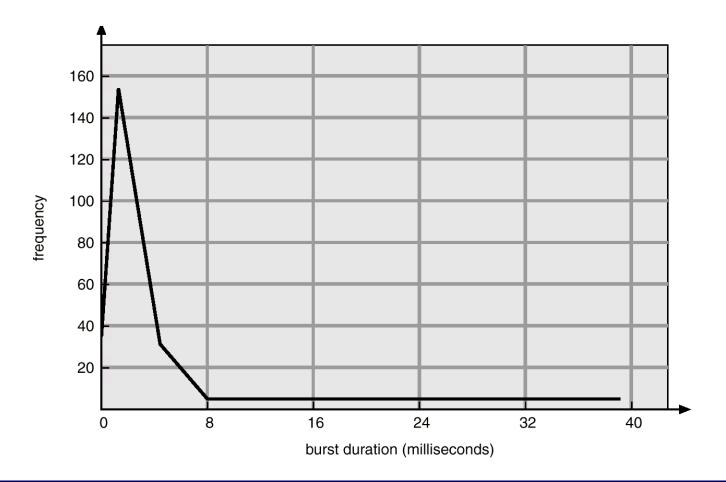
First/last burst: CPU burst

- Types of processes
 - I/O-bound process
 - Consists of many short CPU bursts
 - CPU-bound process
 - Consists of a few long CPU bursts



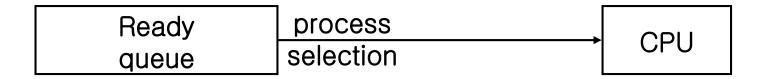
Histogram of CPU-burst Durations

Exponential or hyper exponential



CPU Scheduler

- CPU scheduler (= short term scheduler)
 - Selects a process from ready queue and allocate CPU to it
- Implementation of ready queue
 - FIFO, **priority queue**, tree, unordered linked list
 - Each process is represented by PCB

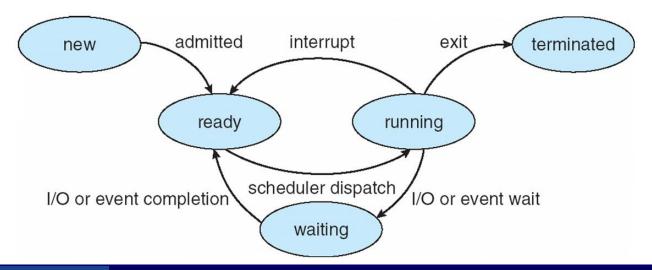


Preemptive Scheduling

- Preemptive scheduling: scheduling may occur while a process is running.
 - Ex) interrupt, process with higher priority

When Scheduling Occurs?

- 1. A process switches from running state to waiting state
- 2. A process switches from running state to ready state ex) an interrupt occurs
- 3. A process switches from waiting state to ready state ex) completion of I/O
- 4. A process terminates
- -> 1 and 4 are inevitable, but 2 and 3 are optional



Preemptive Scheduling



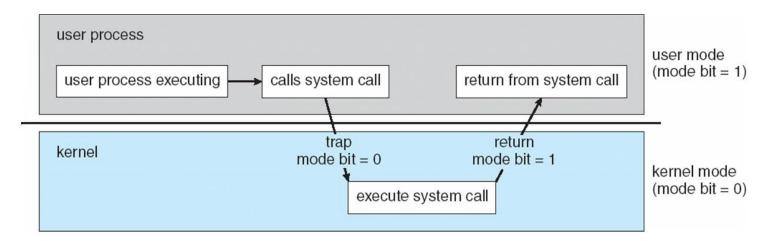
- Scheduling can occur at 1, 4 only
- Running process is not interrupted.

Preemptive scheduling

- Scheduling can occurs at 1, 2, 3, 4
- Scheduling may occur while a process is running
- Requires H/W support and shared data handling

Preemption of OS Kernel

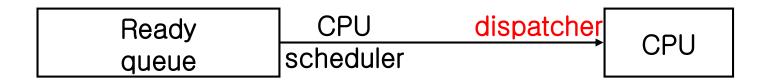
- Preemptive kernel: kernel allows preemption in system call.
 - Desirable for real-time support.



- cf. System calls are non-preemptive in most OS.
 - Keeps kernel structure simple

Dispatcher

- Dispatcher: a module that gives control of CPU to the process selected by short-term scheduler
 - Switching context
 - Switching to user mode
 - Jumping to proper location in user program
- Dispatch latency: time from stopping one process to starting another process



Scheduling Criteria

- CPU utilization: keep the CPU as busy as possible
- Throughput: # of processes completed per time unit
- Turnaroundtime: interval from submission of a process to its completion
- Waiting time: sum of periods spent waiting in ready queue
- Response time: time from submission of a request to first response
- -> Importance of each criterion vary with systems
- Measure to optimize
 - Average / minimum / maximum value
 - Variance

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

Scheduling Algorithms

- First-come, first-served (FCFS) scheduling
- Shortest-job-first (SJF) scheduling
- Priority scheduling
- Round-robin scheduling
- Multilevel queue scheduling
- Multiple feedback-queue scheduling

First-Come, First-Served Scheduling



- Non-preemptive scheduling
- Simplest scheduling method

Sometimes average waiting time is quite long

CPU, I/O utilities are inefficient

Ex) Three processes arrived at time 0

Process	Burst Time	Waiting time
P1	24	0
P2	3	24
P3	2	27

*time unit: msec

Average waiting time:
$$(0 + 24 + 27)/3 = 17$$

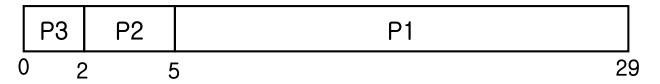
Shortest-Job-First Scheduling

 Assign to the process with the smallest next CPU burst

Process	Burst Time	Waiting time
P1	24	5
P2	3	2
P3	2	0

Average waiting time:

$$(5 + 2 + 0)/3 = 2.3$$



- SJF algorithm is optimal in minimum waiting time
- Problem: difficult to know length of next CPU burst

Shortest-Job-First Scheduling

- Predicting next CPU burst from history
 - Exponential averaging

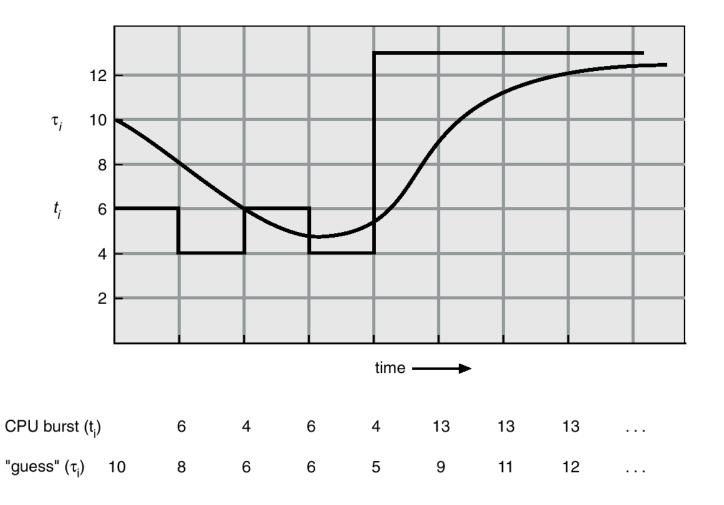
$$\tau_{n+1} = \alpha t_n + (1-\alpha) \tau_n$$

- □ t_n: actual length of n-th CPU burst
- \Box τ_n : prediction for n-th CPU burst
- $\square \alpha$: a coefficient between 0 and 1
 - α = 0: recent history has no effect
 - α = 1: only recent history matters
 - \square Usually, $\alpha = 0.5$

Note: older history affects less

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + ... + (1 - \alpha)^j \alpha t_{n-j} + ... + (1 - \alpha)^{n+1} \tau_0$$

Prediction of CPU Burst



Shortest-Job-First Scheduling



Shortest-remaining-time-first scheduling

Process	Arrival Time	Burst Time	Waiting Time
P1	0	7	9
P2	2	4	1
P3	4	1	0
P4	5	4	2

P1	P2	P3	P2	P4	P1	
(2)	(2)	(1)		(4)	(5)	
0	2	4 5		7	11	16

Average waiting time: 3

Priority Scheduling



Equal-priority processes: FCFS

Note: each process has its priority

In this text, lower number means higher priority

Ex)

Process	Burst Time	Priority	Waiting Time
P1	10	3	6
P2	1	1	0
P3	2	4	16
P4	1	5	18
P5	5	2	1


Average waiting time: 8.2

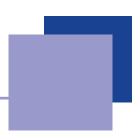
	P2	P5	P1	P3	P4	
0	1		3	6 1	8 1	9

Priority Scheduling

- Priority can be assigned internally and externally
 - Internally: determined by measurable quantity or qualities
 - □ Time limit, memory requirement, # of open files, ratio of I/O burst and CPU burst, ···
 - Externally: importance, political factors
- Priority scheduling can be preemptive or nonpreemptive.
- Major problems
 - Indefinite blocking (= starvation) of processes with lower priorities
 - -> Solution: aging (gradually increase priority of processes waiting for long time)

- Similar to FCFS, but it's preemptive
 - Designed for time-sharing systems
 - CPU time is divided into time quantum (or time slice)
 - □ A time quantum is 10~100 msec.
 - Cf. switching latency: 10 µsec
 - Ready queue is treaded as circular queue
 - CPU scheduler goes around the ready queue and allocate
 CPU time up to 1 time quantum

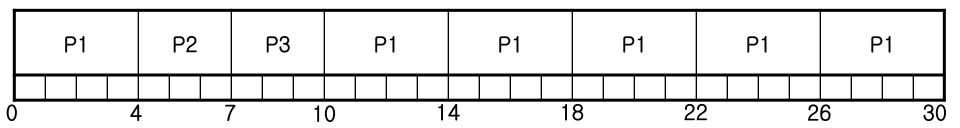




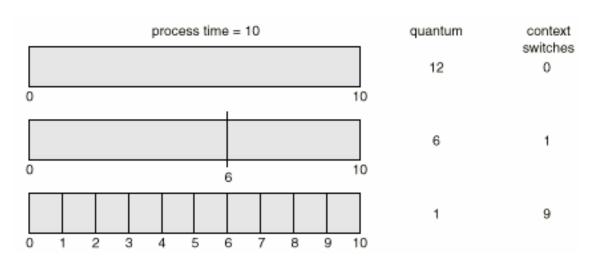
Ex) Time quantum = 4

Process	Burst Time	Waiting Time
P1	24	6
P2	3	4
P3	3	7


Average waiting time: 5.66



- Performance of RR scheduling heavily depends on size of time quantum.
 - Time quantum is small: processor sharing
 - Time quantum is large: FCFS
- Context switching overhead depends on size of time quantum.



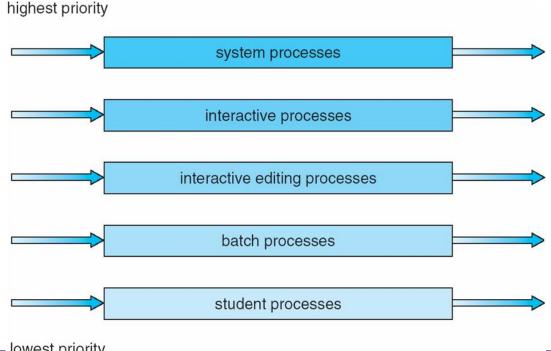
- Turnaroundtime also depends on size of time quantum
 - Average turnaround time is not proportional nor inverse-proportional to size of time quantum
 - Average turnaroundtime is improved if most processes finish their next CPU burst in a single time quantum
 - However, too long time quantum is not desirable
 - A rule of thumb: about 80% of CPU burst should be shorter than time quantum



process	time
P ₁ P ₂ P ₃ P ₄	6 3 1 7

Multilevel Queue Scheduling

- Classify processes into different groups and apply different scheduling
 - Memory requirement, priority, process type, ...
- Partition ready queue into several separate queues



Multilevel Queue Scheduling

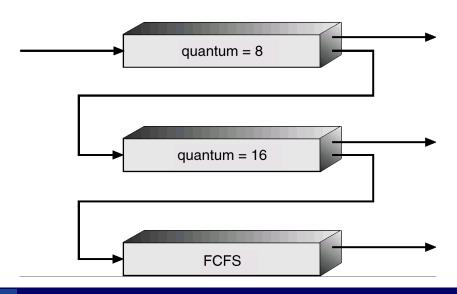
- Each queue has its own scheduling algorithm
- Scheduling among queues
 - Fixed-priority preemptive scheduling
 - A process in lower priority queue can run only when all of higher priority queues all empty
 - Time-slicing among queues
 - Ex) foreground queue (interactive processes): 80% background queue (batch processes): 20%
- Assignment of a queue to a process is permanent.

Multilevel Feedback-Queue Scheduling

- Similar to multilevel queue scheduling, but a process can move between queues.
- Idea: separate processes according to characteristics of their CPU bursts.
 - If a process uses too much CPU time, move it to lower priority queue.
 - I/O-bound, interactive processes are in higher priority queues.

Multilevel Feedback-Queue Scheduling

- Ex) Ready queue consists of three queues (0~2)
 - Q_0 (time limit = 8 milliseconds)
 - Q_1 (time limit = 16 milliseconds)
 - Q_2 (FCFS)
- → A new process is put in Q₀. if it exceeds time limit, it moves to lower priority queue



Multilevel Feedback-Queue Scheduling

- Parameters to define a multilevel feedback-queue scheduler
 - # of queues
 - Scheduling algorithm for each queue
 - Method to determine when to upgrade a process to higher priority queue
 - Method to determine when to demote a process to lower priority queue
 - Method to determine which queue a process will enter when it needs service
- → The most complex algorithm

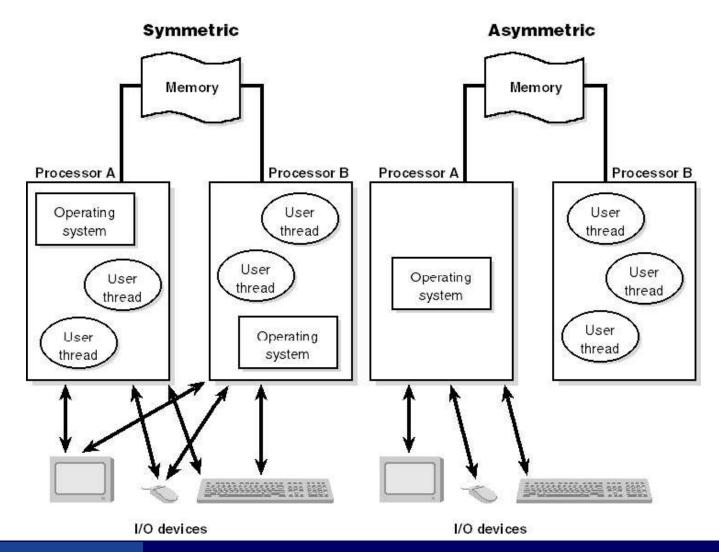
Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

Multiple-Processor Scheduling

- Multiple-processor system
 - Load sharing is possible.
 - Scheduling problem is more complex.
- There are many trials in multiple-processor scheduling, but no generally best solution
- In this text, all processors are assumed identical.
 - Any process can run on any processor

Symmetric vs. Asymmetric Multiprocessing

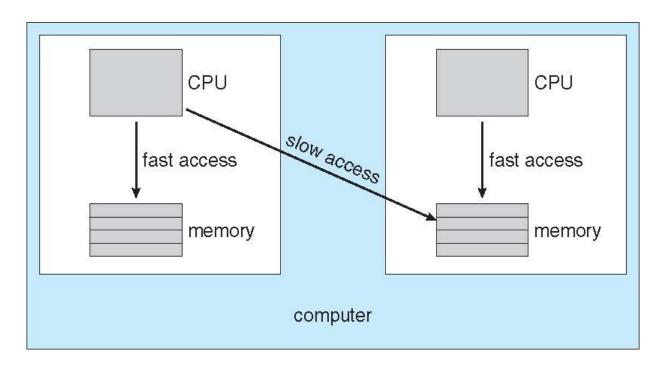


Processor Affinity

- Overhead of migration of processes from one processor to another processor
 - All contents of cache should be invalidated and repopulated
- Processor affinity: keeping a process running on the same processor to avoid migration overhead
 - **Soft affinity**: Although OS attempts to keep a process running on the same processor, a process may migrate between processors.
 - Hard affinity: It is possible to prevent a process from being migrated to other processor

Processor Affinity

- NUMA(Non-Uniform Memory Access) and CPU scheduling
 - A CPU has faster access to some parts of main memory than to other parts



Load Balancing

- Load balancing: attempt to keep the workload evenly distributed across all processors.
 - Necessary for system where each processor has its own ready queue (It's true for most modern OS's)
- Two general approaches
 - Push migration
 - A specific task periodically checks the load on each processor
 - ☐ If unbalance is found, move processes to idle or less-busy processors
 - Pull migration
 - An idle processor pulls a waiting task

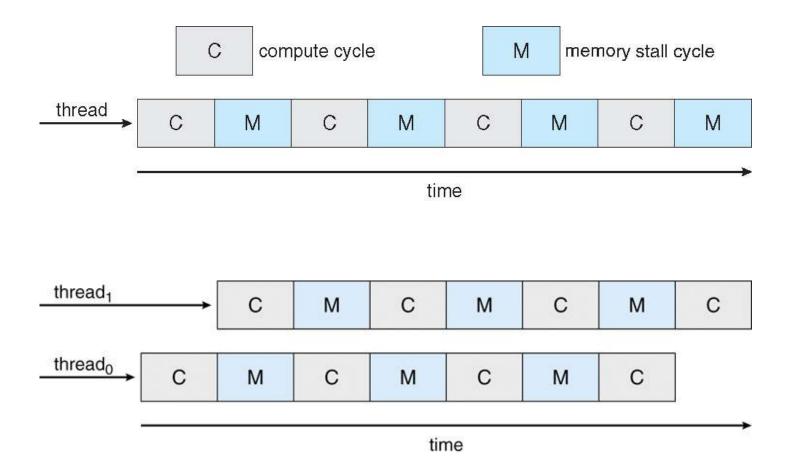
Note! Push and pull migration can be implemented in parallel.

- -> Linux, ULE scheduler for FreeBSD
- Load balancing can counteract the benefit of processor affinity

Multi-core Processors

- Multi-core processor: multiple processor cores on the same physical chip
 - Recent trend
 - Faster and consume less power than systems in which each processors has its own physical chip
- Scheduling issues on multi-core processor
 - Memory stall: for various reasons, memory access spends significant amount time waiting for the data to become available.
 - Ex) accessing data not in cache
 - Remedy: multithreaded processor cores
 - □ Two or more hardware threads are assigned to each core

Multi-Threaded Processor Cores

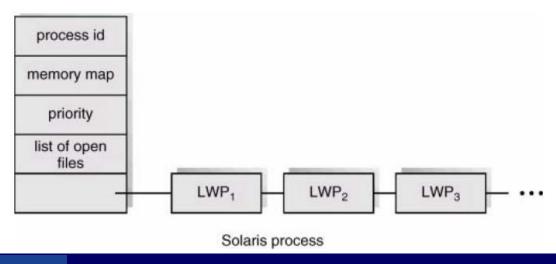


Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

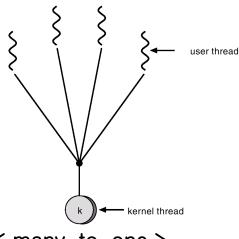
Thread Scheduling

- Threads
 - User thread -> supported by thread library
 - Scheduled indirectly through LWP
 - Kernel thread -> supported by OS kernel
- Actually, it is not processes but kernel threads that are being scheduled by OS

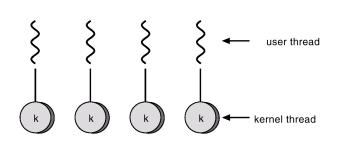


Contention Scope

- Process-contention scope (PCS)
 - Competition for LWP among user threads
 - Many-to-one model or many-to-many model
 - Priority based
- System-contention scope (SCS)
 - Competition for CPU among kernel threads




< many-to-one >



< one-to-one >

Scheduler Activation and LWP

Connection between user / kernel threads through LWP



User-level thread



L Light-weight Process



Pthread Scheduling

- In thread creation with Pthreads, we can specify PCS or SCS
 - PCS(PTHREAD_SCOPE_PROCESS):
 - □ In many-to-one, only PCS is possible
 - SCS(PTHREAD_SCOPE_SYSTEM):
 - □ In one-to-one model, only SCS is possible

Note: On certain system, only certain values are allowed

- □ Linux, MacOS X
- Related functions
 - pthread_attr_setscope(pthread_attr_t *attr, int scope)
 - pthread_attr_getscope(pthread_attr_t *attr, int *scope)

Pthread Scheduling

Example)

```
pthread t tid;
pthread_attr_t attr;
int scope;
pthread_attr_init(&attr);
pthread_attr_getscope(&attr, &scope);
// scope is either PTHREAD_SCOPE_PROCESS
// or PTHREAD_SCOPE_SYSTEM
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM)
pthread_create(&tid, &attr, runner, NULL);
```

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

Operating System Examples

Solaris

Windows XP

Linux

Windows Scheduling

- Priority-based, preemptive scheduling
 - Dispatcher handles scheduling
 - Dispatcher ensures the highest priority thread will always run
 - 32-level priority scheme
 - □ Variable class (1~15)
 - □ Real-time class (16~32)
 - Memory management (0)
 - If no ready thread is found, dispatcher runs idle thread

Windows Scheduling

Priority classes

- A process belongs to one of 6 levels of priority class
- Each thread belongs to one of 7 levels of relative priority

	real- time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Windows Scheduling

Altering priority

- Priority class of a process: SetPriorityClass()
- Base priority of a thread: SetThreadPriority()

Adjusting priority for variable-priority class

- Lower priority of a thread that expired entire time quantum
 - □ Never lowered below the base priority
- Boost priority of a thread that released from waiting
 - Amount of increment vary with what it was waiting

Scheduling of foreground / background processes

 Time quantum for the foreground process is increased by some factor (typically, 3)

Linux Scheduling

History

- Two problems in Linux prior to version 2.5
 - Not adequate support for SMP system
 - Not scale well as # of tasks increases
- Overhauled scheduler in Linux version 2.5
 - □ Scheduling complexity is O(1)
 - □ Increased support for SMP
 - □ Poor response time interactive processes
- Linux 2.6
 - Completely Fair Scheduler (CFS)

Linux CFS Scheduler

Based on scheduling classes

- Each class is assigned a specific priority.
 - □ Default scheduling class
 - □ Real-time scheduling class
 - □ If necessary, new scheduling classes can be added
- Different scheduling algorithms for different scheduling classes
- Scheduler selects highest-priority task belonging to highest-priority class

Relative priority

- 'Nice value' from -20 to +19
 - ☐ Higher value represents lower priority ('nice' to others)

Linux CFS Scheduler

- CFS scheduler assigns a proportion of CPU processing time, rather than discrete values of time slices, to each task
 - The portion is calculated from nice value.
 - Proportions of CPU time are allocated from the value of targeted latency.
 - □ Targeted latency: interval of time during which every runnable task should run at least once
 - Targeted latency can increase if the number of active tasks in the system grows beyond a certain threshold

Linux CFS Scheduler

- Per task variable vruntime(virtual run time) measures how long each task has run
- vruntime of each task is associated with a decay factor based on the priority of a task
- Higher-priority task can preempt a lower-priority task.

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Operation system examples
- Algorithm evaluation

- First of all, evaluation criteria should be determined Ex)
 - 1. Maximum response time is 1 second
 - 2. Satisfying 1, maximize CPU utilization
 - 3. Maximize throughput such that turnaround time is linearly proportional to total execution time

Evaluation methods

- Deterministic modeling
- Queueing models
- Simulations

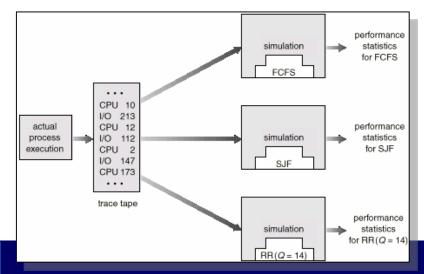
Deterministic modeling

Evaluation for a particular predetermined workload

Queueing models

- Evaluation based on probabilistic distribution of...
 - CPU and I/O burst
 - Arrival time of CPU burst
- Description of system
 - □ Network of servers, each of which has a ready queue
 - Given arrival rates and service rates, performance can be computed probabilistically

- Simulations, using a programming model of computer system
 - Generator of process, CPU burst time, arrival, departure, …
 - Usually, the generator randomly generates the simulated events based of some distribution
 - Limitation: simulated situation isn't exactly same with the real situation
 - □ Especially, distribution gives no information about order of events
 - Remedy: using a record of real system. (trace tape)



- Implementation: the only way for accurate evaluation
 - Performance depend not only on scheduling algorithm and OS support, but also user's interaction
 - Environment will change
 - Ex) If short process is given higher priority, user may break a large process into several smaller processes.
 - Some OS's provides flexibility to alter scheduling scheme
 - Performance turning for specific (set of) applications
 - API's to modify priority of a process or thread