

7. Deadlocks

ECE30021/ITP30002 Operating Systems


Agenda



- Introduction
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection
- Recovery from deadlock

System Model


- A system consists of ...
 - A finite number of resources
 - Resources are partitioned into several **types**
Ex) two identical printers
 - A number of competing processes



System Model

■ Steps to utilize a resource

1. Request
 - Process should wait until the request is granted
2. Use
3. Release



Deadlock State

- A set of processes is in deadlock state
 - Each process is waiting for an event that can be caused only by another process in the set.
 - Mainly concerned with resource acquisition and release



Deadlock Characterization



- **Necessary conditions for deadlock**
 1. Mutual exclusion
 2. Hold and wait
 3. No preemption
 4. Circular wait

Resource-Allocation Graph


■ $G(V, E)$

■ V : two types of vertices

- $P = \{P_1, P_2, \dots, P_n\}$: a set of processes
- $R = \{R_1, R_2, \dots, R_m\}$: a set of resource types

■ E : two types of directed edges



- $P_i \rightarrow R_j$: **request edge**
- $R_j \rightarrow P_i$: **assignment edge**
- If a request edge is fulfilled, it becomes an assignment edge



Resource-Allocation Graph

■ Cycle vs. deadlock state

- No cycle \rightarrow no deadlock
- Cycle \rightarrow deadlock (?)
 - Only one instance per resource type – deadlock occurred
 - Several instances per resource type – possibility of deadlock



Methods for Handling Deadlock



- Three methods to handle deadlock
 - Prevent or avoid deadlock
 - Detect deadlock and recover
 - Ignore
 - Most common case
 - Deadlock is very infrequent (once a year)

Methods for Handling Deadlock



■ Deadlock prevention

- A set of method for ensuring at least one of necessary conditions cannot hold
- Constraint on request for resources

■ Deadlock avoidance

- Keep the system in **safe state** in which deadlock cannot occur using additional information, such as ...
 - Resources currently available or allocated to each process
 - **Additional information about future requests and release of each process.**

Agenda



- Introduction
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection
- Recovery from deadlock

Deadlock Prevention



- **Mutual exclusion**
 - Possible only for sharable resources.
 - Ex) read-only file
 - Many resources are intrinsically non-sharable.

- **Hold and wait: guarantee whenever a process requests a resource, it does not hold other resources**
 - Allocate all required resources before it begins execution.
 - Allow a process to request resources only when it has none.

→ Inefficient utilization of resource

→ Starvation

Deadlock Prevention



■ No Preemption

- If a process cannot allocate a resource immediately, it should release all resources it hold.
- The process restarts when it regains its old resources and the new resource it requested.

■ Circular wait

- Impose total ordering of all resource types
- All resources should be requested in that order.

Ex) $F(\text{tape drive}) = 1$

$F(\text{disk drive}) = 5$

$F(\text{printer}) = 12$

Deadlock Prevention

■ Deadlock

P_0	P_1
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
\vdots	\vdots
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q);</i>	<i>signal(S);</i>

Deadlock Prevention



- Problems of deadlock prevention
 - Low device utilization
 - Reduce system throughput

Deadlock Avoidance



- Require **additional information** about how resources are to be requested.
 - Ex) Each process declares maximum # of resources of each type it may request.
- Deadlock avoidance algorithm dynamically examines **resource-allocation state** to ensure a circular-wait condition can never exist.
- Resource-allocation state is defined by
 - # of available resources
 - # of allocated resources
 - **Maximum demand of processes**



Safe State



- When a process requests an available resource, system must decide if immediate allocation leaves the system in **safe state**.
- **Safe sequence**: a sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resource requests of P_i can be satisfied by ...
 - Currently available resources, and
 - Resources held by P_j , where $j < i$
- **Safe state**: there exists a **safe sequence** of all process


Safe Sequence

- Are they safe ?



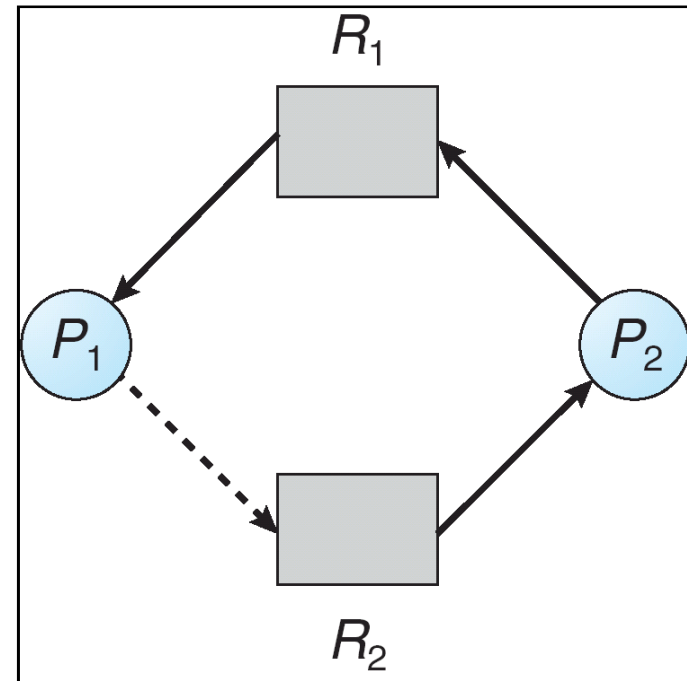
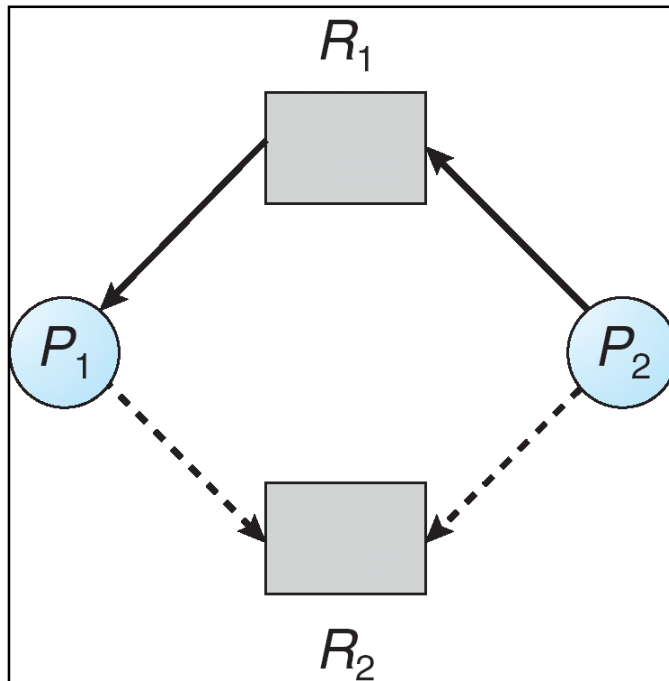
Safe State and Deadlock

- If a system is in safe state → no deadlock
- If a system is in unsafe state → possibility of deadlock
- Deadlock avoidance → keeping a system in safe state



Detecting Unsafe State

- Claim edge $P_i \rightarrow R_j$ (dashed line): process P_i **may request** resource R_j at some time
 - Claim edge can be converted into request edge



unsafe state !

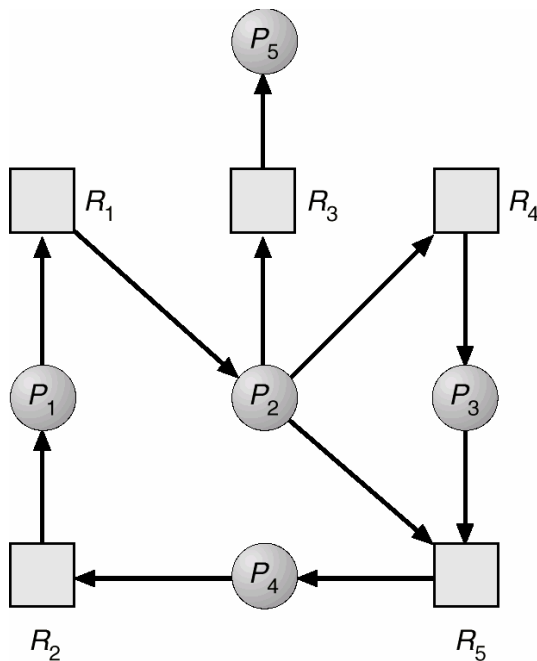
Deadlock Detection



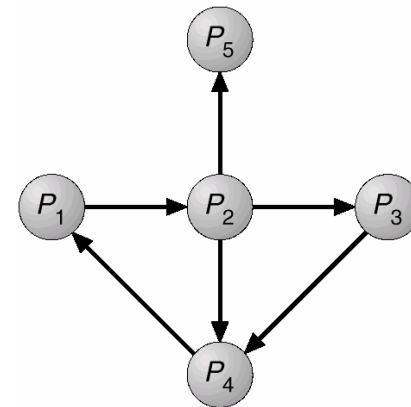

- Single instance of each resource type
- Several instances of a resource type (→ skip)

Single Instance of Each Resource Type

- **Wait-for graph**: removing resource nodes from resource-allocation graph and collapsing the appropriate edges



Resource-allocation graph



Wait-for graph

Single Instance of Each Resource Type



- A deadlock exists if and only if the wait-for graph contains a cycle.
- To detect deadlock, system should
 - Maintain wait-for graph
 - Invoke an algorithm to detect a cycle – $O(n^2)$

Detection–Algorithm Usage



- When should we invoke the detection algorithm?
 - How often is a deadlock likely to occur?
 - How many processes will be affected by deadlock when it happens?

- An observation
 - Deadlocks occur only when some processes makes a request that cannot be granted immediately
 - We may invoke detection algorithm whenever a request for allocation cannot be granted immediately
(In this case, we can find the process which caused deadlock.)

Recovery from Deadlock



- Process termination
- Resource preemption

Process Termination



- Abort all deadlocked processes
 - Too expensive
- Abort one process at a time until the deadlock is eliminated
 - Order of priority
 - Time from start / time to completion
 - Resources the process has used / needs to complete
 - How many processes will need to be terminated?
 - Is the process interactive or batch?

Resource Preemption



■ Selecting a victim

- Minimizing cost (# of resources a process has, amount of time consumed so far, ...)

■ Rollback

- The selected process should return to some safe state and restart it.

■ Starvation

- How can we guarantee that resources will not always be preempted from the same process?