# 8. Main Memory

ECE30021/ITP30002 Operating Systems

### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

#### Background

#### Importance of memory management

- In multitasking system, several processes may exist in memory
- Memory is central of modern computer system

	<u>Main memory</u>				
0				•••	
4 8				•••	
8				•••	
12 16				•••	
16				•••	
•••	•••				

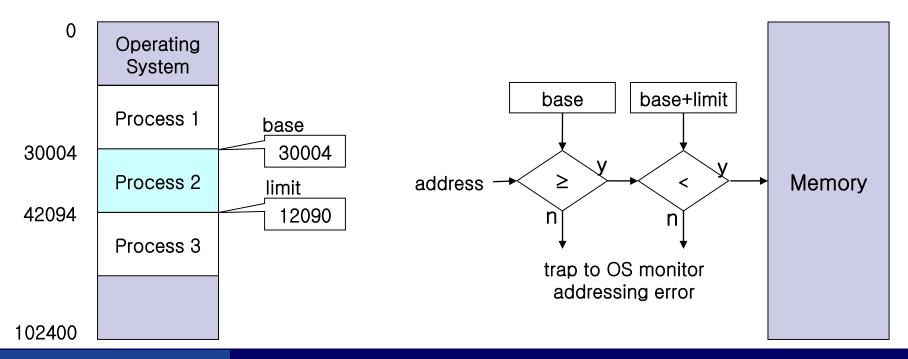
#### Memory structure

- A large array of bytes each with its own address.
- Memory unit is interested only in the sequence of memory addresses.

#### **Basic Hardware**

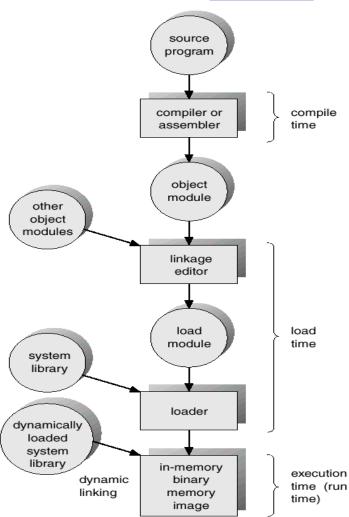
#### Protection from other processes

- Each process has a separate memory space.
- Can be implemented by base register and limit register.
- CPU H/W compares every address generated.
- Base/limit registers can be loaded by privileged instruction.



#### **Address Binding**

- Most systems allow a user process to reside in any part of the physical memory.
- Representation of address
  - Source program: symbolic address Ex) count
  - Compiler
    - Symbol -> relocatable address
       Ex) 14 bytes from beginning of the module
  - Linkage editor or loader
    - Relocatable address -> absolute address
       Ex) 0x74014
- Address binding: mapping one address space to another



# **Address Binding**



#### **Address Binding**

#### Address binding time

- Compile time: absolute code
  - □ MS-DOS .com file
  - □ If starting location changes, the program should be recompiled
- Load time: relocatable code
  - If starting location changes, the program should be reloaded.
- Execution time
  - Ex) Systems that support logical addressing
  - □ Used by most general-purpose OS's
  - Requires H/W support for address mapping

## Logical vs. Physical Address Space

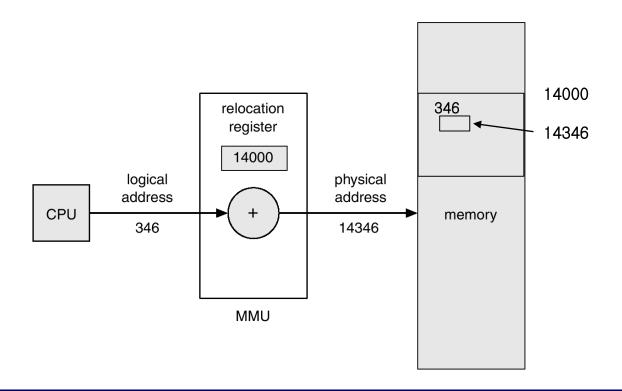
- Logical address: address generated by CPU in common cases
- Physical address: address seen by memory unit

### Logical vs. Physical Address Space

- Logical address vs. physical address
  - In compile-time and load-time address binding
    - □ Logical address = physical address
  - In execution-time address binding
    - □ Logical (virtual) address ≠ physical address
- Address spaces
  - Logical address space: set of all logical addresses generated by a program
  - Physical address space: set of all physical addresses corresponding to these logical addresses
- Run-time address mapping is done by memory-management unit (MMU)

### Logical vs. Physical Address Space

- A simple method of address mapping using relocation register (base register)
  - User program never sees physical addresses



## Dynamic Linking and Shared Library

- Static linking: all object modules are combined into the binary program image
- Dynamic linking: linking is postponed until execution time Ex) System libraries
  - Stub: a small piece of code that indicates …
    - How to locate the appropriate library routine
    - □ How to load the library if it is not ready
  - After a library routine is dynamically linked, that routine is executed directly
    - stub replaces itself with address of the routine and execute it.
- Advantages of dynamic linking
  - Library code can be shared among processes.
    - Multiple versions of a library can coexist
  - Library update doesn't require re-linking.
- Dynamic linking requires help from OS.

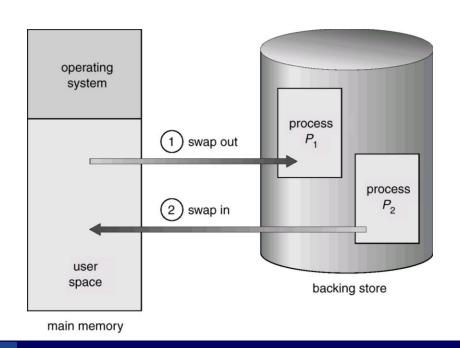


### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

#### Swapping

- When a process starts, it must be in memory for execution.
- However, a process can be swapped to a backing store and back into memory to continue execution.



#### Swapping

- A process swapped-out will be swapped back into the same memory space it occupied previously
  - Exception: execution-time binding
- Backing store: fast disk (usually, separated from file system)
  - Should be large enough to accommodate copies of all memory images for all users
  - Must provide direct access to memory images
- Dispatching
  - Firstly see if the process selected by scheduler is in memory.
  - If not, it swaps out a process and swaps in the desired process

#### Swapping

Context-switch time in swapping system

Ex) size of a process: 100 MB disk transfer rate: 50 MB/s average latency: 8 ms

Then, swapping time: (100MB / 50MB/s + 8ms) \* 2 = 4016 ms

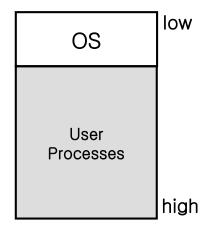
- \* Swapping time is affected by amount of memory swapped
  - □ We would need to swap only what is actually used.
- Time slice should be long relative to swap time
- Many systems use modified versions of swapping
  - Swapping is enabled only when memory usage exceeds some threshold.

### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

## Contiguous Memory Allocation

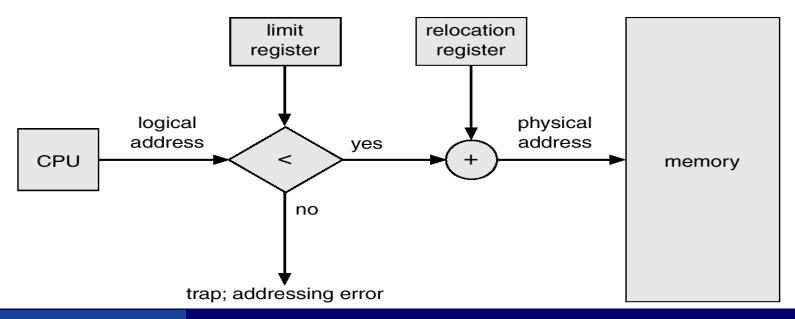
- Memory is usually divided into two partitions
  - Resident OS
    - □ Low memory with interrupt vector in most systems
  - User processes
    - High memory



- Several user processes can reside in memory at the same time
  - -> How to allocate memory to processes waiting in input queue?

### Memory Mapping and Protection

- Memory mapping and protection are provided by relocation register and limit register
  - Add memory access, MMU maps the logical address dynamically.
  - Relocation and limit registers are loaded by dispatcher during context switching.



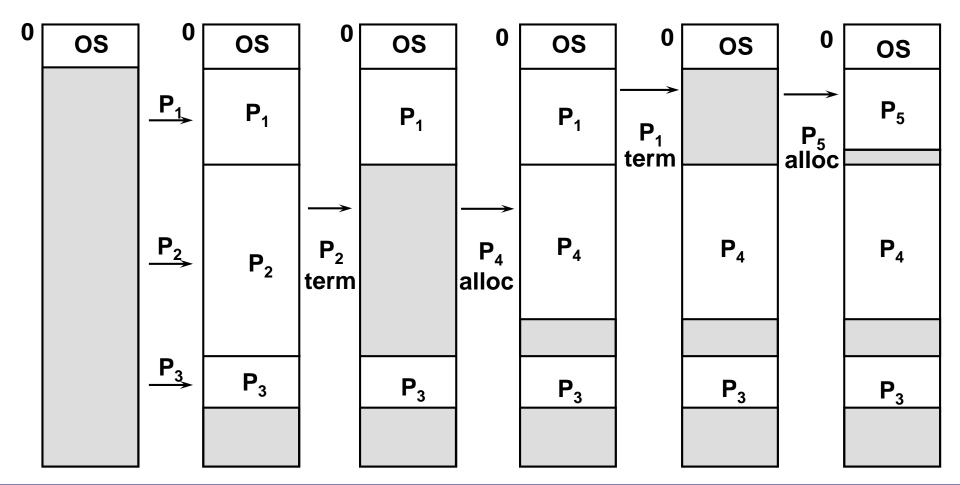
### **Memory Allocation**



- Initially all available memory forms a large block of free memory (hole)
- If a process arrives and needs memory, search a hole large enough for this process.
- If a process terminates, it releases its memory
  - Adjacent holes can be merged
- \* MVT: Multiprogramming with a Variable number of Tasks

### **Memory Allocation**

Variable-partition scheme



#### Fragmentation

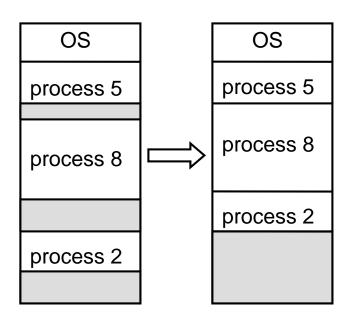
- External fragmentation: free memory space broken into little pieces
  - Sometimes, it's impossible to allocate large continuous memory although total size of free memory larger than the required memory
  - 50-percent rule (analysis of first fit)
    - □ Given N allocated blocks, another 0.5N blocks will be lost to fragmentation. (1/3 of memory may be unusable)
- Fragmentation of main memory also affects backing store

#### Fragmentation

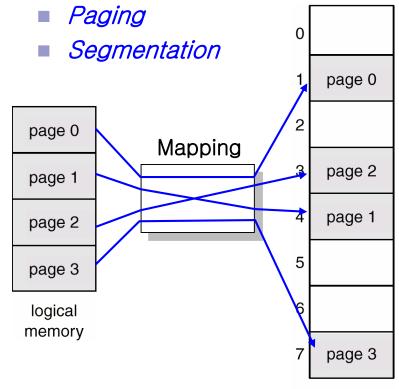
- Memory allocated to a process may slightly larger than requested memory.
  - Ex) allocating 18462 bytes from hole whose size is 18464 bytes
  - Reduce overhead to keep small holes.
  - Generally, physical memory can be broken into fixed-size blocks.
- Internal fragmentation
  - Difference between the amounts of allocated memory and the requested memory.

#### Fragmentation

- Remedy of external fragmentation: compaction
  - Possible only if relocation is dynamic and is done at execution time
  - Expensive



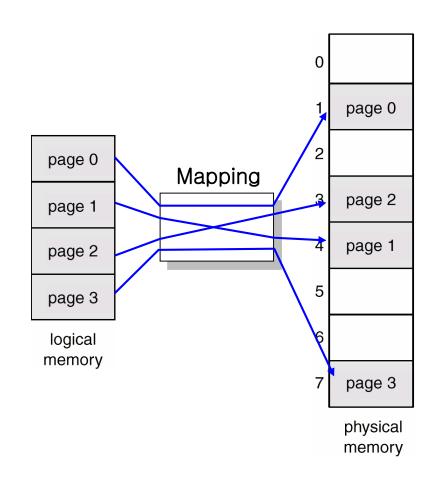
 Alternative remedy: permit logical address space of the processes to be noncontiguous



### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

- A memory-management scheme that permits the physical address space of a process to be noncontiguous.
- Physical memory: consists of frames (fixed-size blocks)
- Logical memory: consists of pages (fixed-size blocks)
- Pages are mapped with frames through page table



Paging model of logical and physical memory

page 0 page 1 page 2 page 3 logical memory

0 page table

frame number 0 page 0 2 page 2 page 1 5 6 page 3 physical

memory



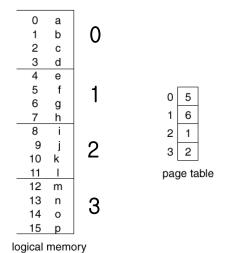
- Size of logical address space: 2<sup>m</sup>
- Page size: 2<sup>n</sup>

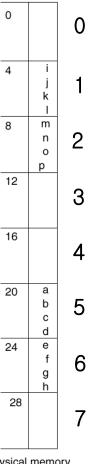
page number	page offset	
p	d	
	n	

page 0
page 1
page 2
page 3

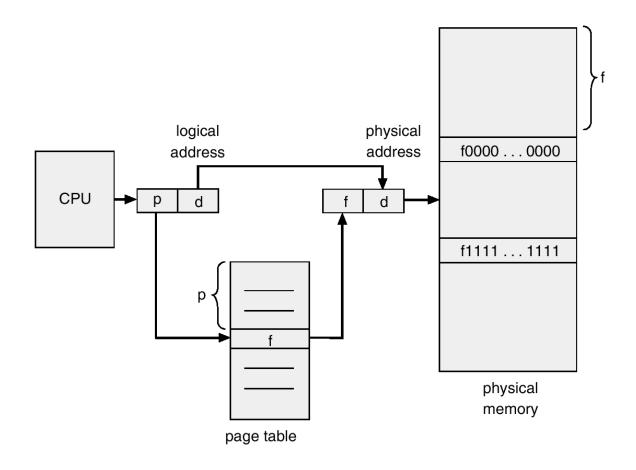
logical memory

#### An example of paging





Paging hardware



#### Overhead of paging

- Page table
- Internal fragmentation
  - □ In average, half page per process

    Note! In paging, no external fragmentation

#### In many systems …

- Page size: 4 KB ~ 4 MB
- Page table entry: 4 bytes (32 bits)

### X86 Page-Table Entry

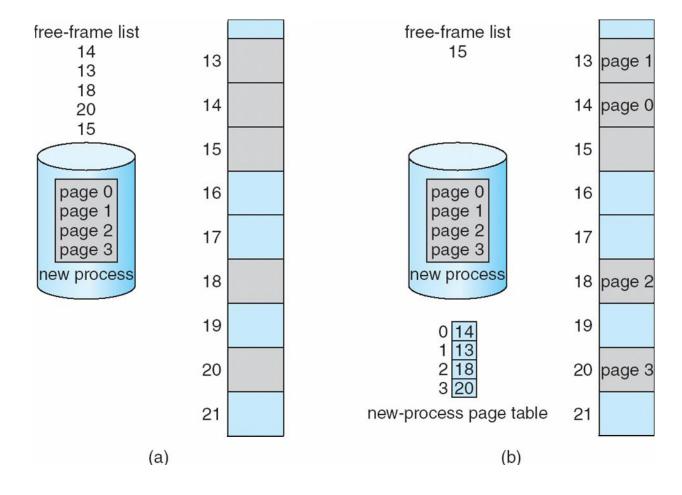
Format of x86 page table entry

#### Page-Table Entry (4-KByte Page) 31 12 11 9876543210 G A D A C W / / P Page Base Address Avail Available for system programmer's use ——— Page Table Attribute Index———— Dirty — Accessed ———————— Cache Disabled ————— Write-Through — Read/Write — Present ————————

#### Frame Table

- OS keeps information about frames in a *frame table* 
  - # of total fames
  - Which frames are allocated
  - Which frames are available
  - ETC.
- OS maintains list of free frames

#### Free Frames



**Before allocation** 

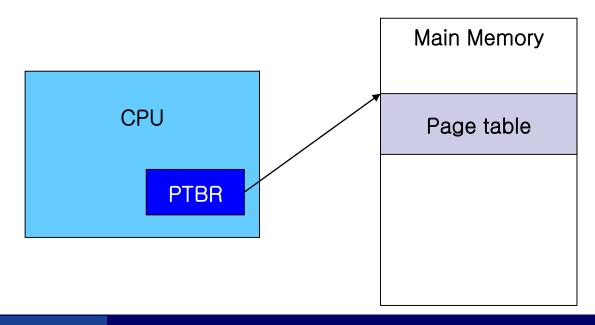
After allocation

### Hardware Support

- Most OS's allocate a page table to each process.
  - Pointer to the page table is stored with register values (in PCB).
  - In context-switching, dispatcher also loads the correct H/W page table.
- H/W implementation of page table
  - A set of dedicated registers
    - Efficiency is very important.
    - Ex) DEC PDP-11: page table (8 entries) is kept in fast registers
    - -> not feasible if page table is large

### Hardware Support

- Page-table base register (PTBR)
  - Page table is stored in main memory
  - PTBR points the page table (eg. CR3 register of x86)
  - cf. Remaining problem: overhead to access page table in main memory (requires two accesses to access a byte.)

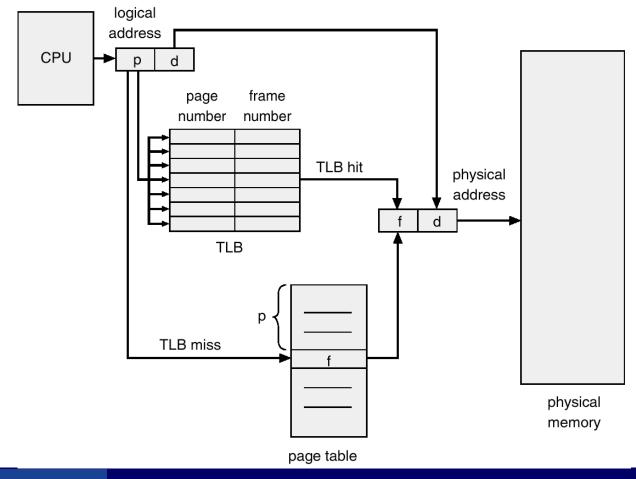


#### Hardware Support

- Translation look-aside buffer (TLB): small fast look up H/W cache.
  - Associative, high-speed memory, consists of pairs (key, value)
    - If page number of logical address is in TLB, delay is less than 10% of unmapped memory access
    - □ Otherwise (TLB miss), access page table in main memory
      - □ Insert (page number, frame number) into TLB
      - ☐ If TLB is full, OS select one for replacement
  - Some TLB's store address-space identifiers (ASID) in each entry (ex: MIPS)
    - ASID identifies process
    - □ Processors w/o ASID (ex: x86) flush entire TLB on process switch

### Hardware Support

#### Paging using TLB



### **Associative Memory**

Associative memory – parallel search

Page #	Frame #	

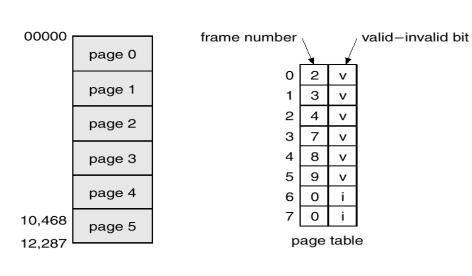
- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory

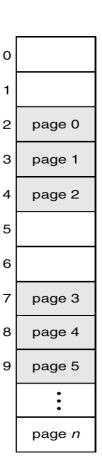
#### **Protection**

- Each frame has a protection bits
  - Read-write / read only
  - Attempt to write to a read-only page causes H/W trap
  - -> Extension: read-only, read-write, execute-only, ...
    - Combination of them
- Each entry of page table has valid-invalid bit
  - OS does not allows to access the page if it is invalid
- Rarely does a process use all of its address range
  - Page table can be wasting memory
  - -> some system provides page-table length register (PTLR)

#### Valid-Invalid Bit

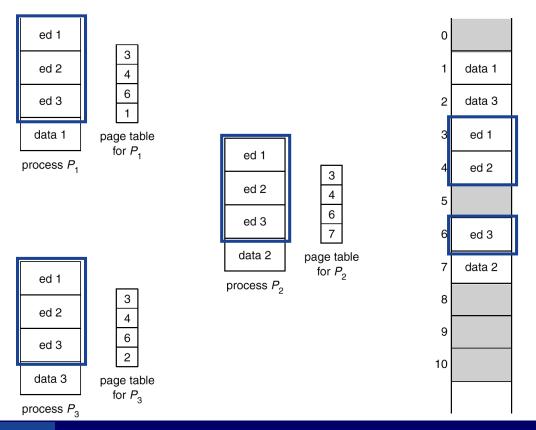
Valid-invalid bit in a page table





#### **Shared Pages**

- Paging provides possibility to share common code
  - Ex) Process P1, P2, P3 runs the same text editor with different data (editor should be non-self-modifying code)



#### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

#### Hierarchical Paging

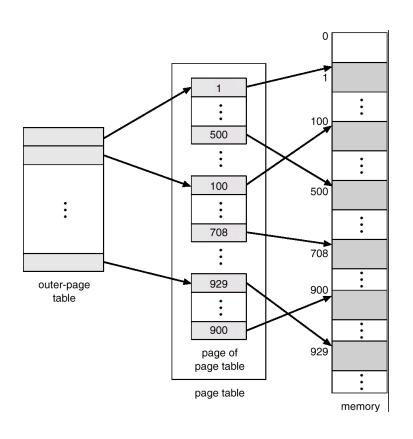
 Motivation: If a system supports large logical address space, page table itself becomes significant overhead

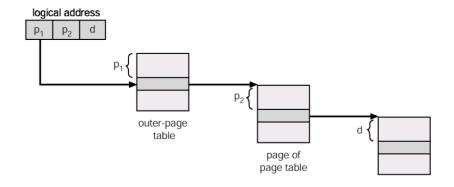
Ex) 32-bit address space, page size is 4 KB( $2^{12}$ ) size of page table =  $2^{(32-12)} > 1$  Million

- Difficult to allocate contiguous memory
- Hierarchical paging
  - Page table itself is also paged
    - Structure of logical address

page number		page offset	
$p_{i}$	$p_2$	d	
10	10	12	

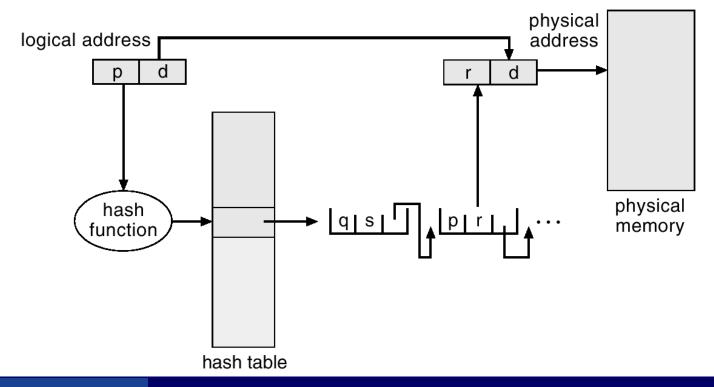
# Hierarchical Paging





#### Hashed Page Table

- Use hashing to for page table
  - Common approach for large address spaces (> 32bits)
  - Each location of hash table consists of linked list of page table entries



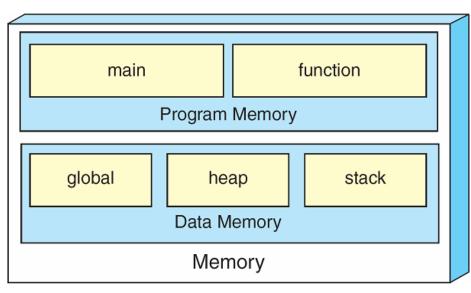
#### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

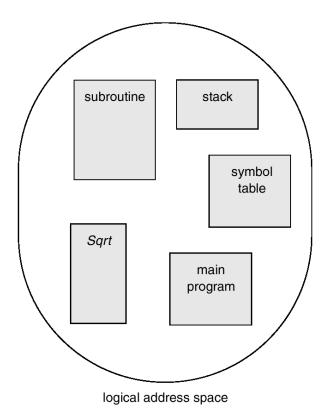
#### Segmentation

 Memory from user's view: collection of variable-size segments, with no necessary ordering

Ex) stack, math library, main program, ...



Conceptual view of memory



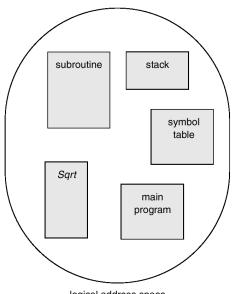
#### Segmentation

- Segmentation: memory-management scheme that supports user view of memory
  - Logical address space is a collection of segments
  - Each segment has a name (number) and length

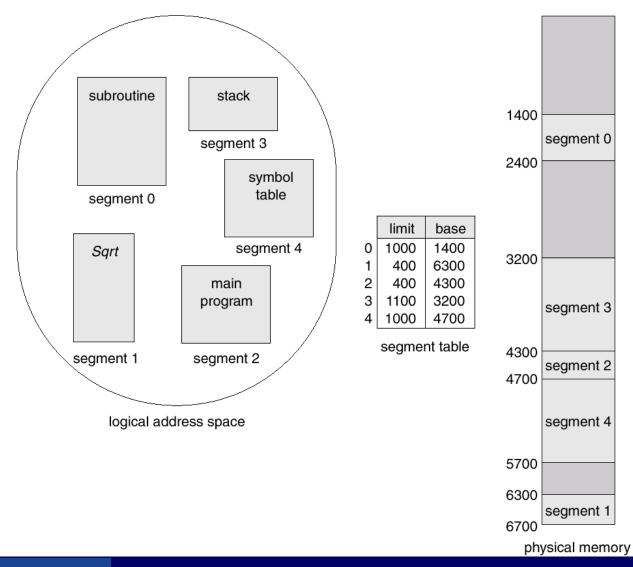
Logical address = <segment-number, offset>

Ex) C compiler might create segments such as ...

- Code
- ☐ Global variables, heap, stacks
- □ Standard C library

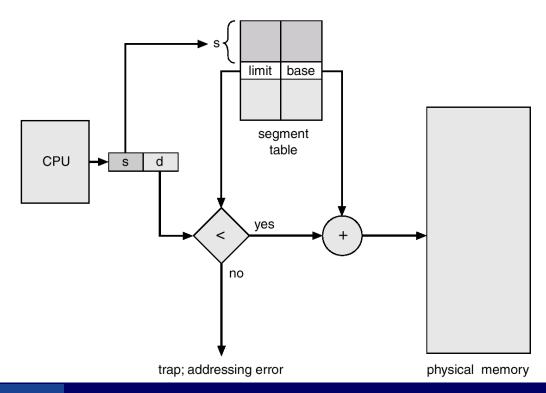


# Example of Segmentation



#### Hardware for Segmentation

- 2D address should be mapped 1D sequence of bytes
- Segment table
  - Consists of segment base and segment limit

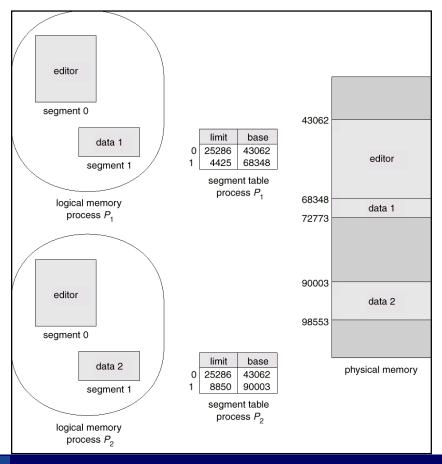


#### Segmentation Architecture

- Segment table maps two-dimensional physical addresses; each table entry has:
  - base contains the starting physical address where the segments reside in memory
  - limit specifies the length of the segment
- Segment-table base register (STBR) points to the segment table's location in memory
- Segment-table length register (STLR) indicates number of segments used by a program; segment number s is legal if s < STLR</p>

#### Segmentation

Protection and sharing is also possible for segmentation

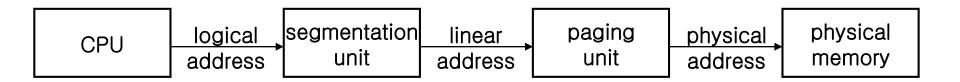


#### Agenda

- Background
- Swapping
- Contiguous memory allocation
- Paging
- Structure of page table
- Segmentation
- Example: Intel Pentium

#### Example: Intel Pentium

- Some architectures supports both paging and segmentation
  - Intel x86: pure segmentation, segmentation with paging



- Maximum size of a segment: 4 GB
- Maximum # of segments: 16 K (=16,384)
  - 1st partition for private segments (8 K)
    - -> kept in local descriptor table (LDT)
  - 2<sup>nd</sup> partition 8 K for shared segments (8 K)
    - -> kept in global descriptor table (GDT)

#### Pentium has ···

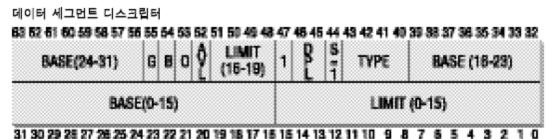
- Six segment registers (CS, SS, DS, ES, FS, GS)
- Six 8-byte microprogram registers to hold descriptors from LDT or GDT
  - Cache for LDT/GDT entry

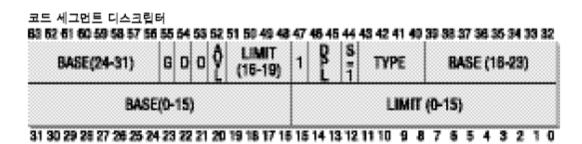


#### Segment descriptor

64 bit information that describes attribute of a segment

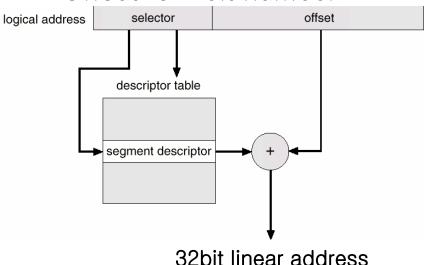
<Structure of
segment descriptors>

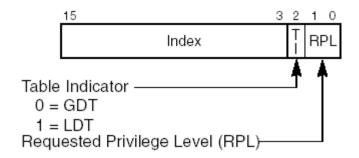


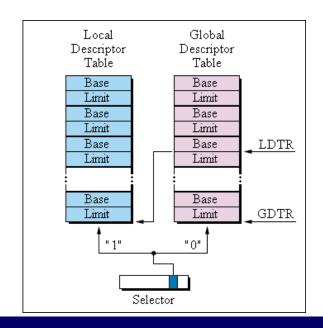


- Logical address: pair of (selector, offset)
  - Selector: 16 bit number
    - □ Segment number (13 bits=8 k)
    - □ GDT or LDT (1 bit)
    - Protection (2 bits)

Offset: 32-bit number







#### Pentium Paging

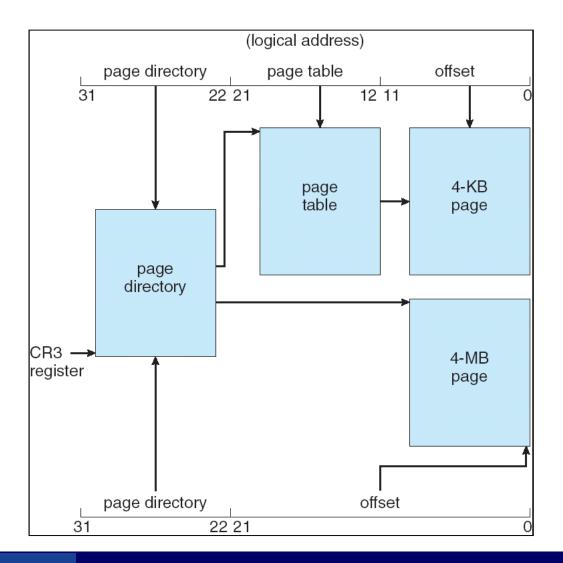
- Pentium supports a page size of 4 KB or 4 MB
  - Page directory entry includes page size flag
  - 4 KB page: two-level paging (page size flag = 0)
    - □ PSE bit of CR4 register

page number		page offset	
p <sub>1</sub> (10 bits)	p <sub>2</sub> (10 bits)	d (12 bits)	

4 MB page (page size flag = 1)

page number	page offset
p <sub>1</sub> (10 bits)	d (22 bits)

# Pentium Paging



### Linux on Pentium Systems

- Linux does not rely on Pentium segmentation but uses it minimally
- Linux on Pentium uses six segments
  - Kernel code / data
  - User code / data
  - Task-state segment (TSS)
  - Default LDT segment
    - Shared for all processes, but usually not used
- Segments for user code and user data are shared by all processes
  - All processes share the same logical address space
  - All segment descriptors are stored in GDT
    - □ If necessary, a process can create LDT and use it

### Linux on Pentium Systems

- Linux uses three-level paging model
  - Linear address in Linux

global middle directory	page table	offset
-------------------------	---------------	--------

- -> works well for both 32 bit and 64-bit architectures
- On Pentium, size of middle directory is zero bit